

## WELCOME TO THE WORLD OF EXTREME

At Leanfra we deliver the highest performance computing with a price that fits everyone!

Access our environment for free

Full Name

Email

Phone

Company name/Job title

SEND REQUEST →



### NETWORKING →

200Gb/s Infiniband  
Lightspeed fast



### GPU SERVERS →

Multi-tiered, latest generation GPUs powered by NVIDIA®



### STORAGE →

Robust. Scalable. Best in class throughput and IOPs

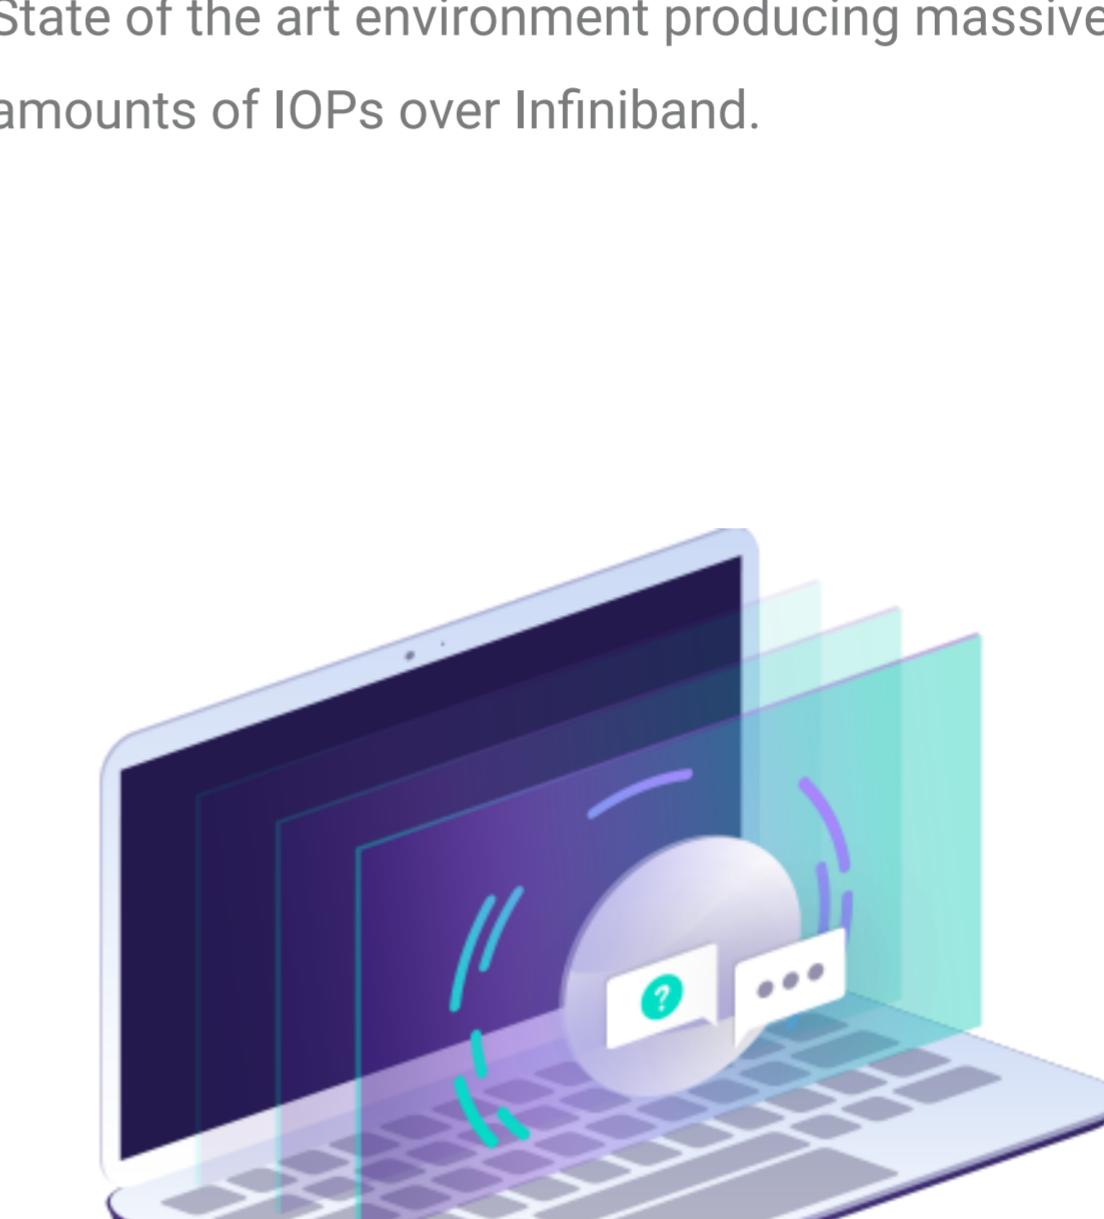
## Why Leanfra?

Evolving infrastructure to the extreme



### Bare-metal freedom

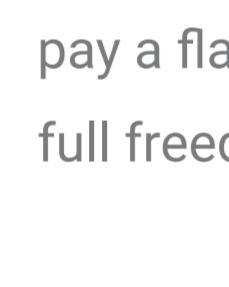
Raw high performance bare metal computing, several tiers to choose from. High end enterprise grade hardware.



→  
→  
→  
**200Gb/s**  
→  
→

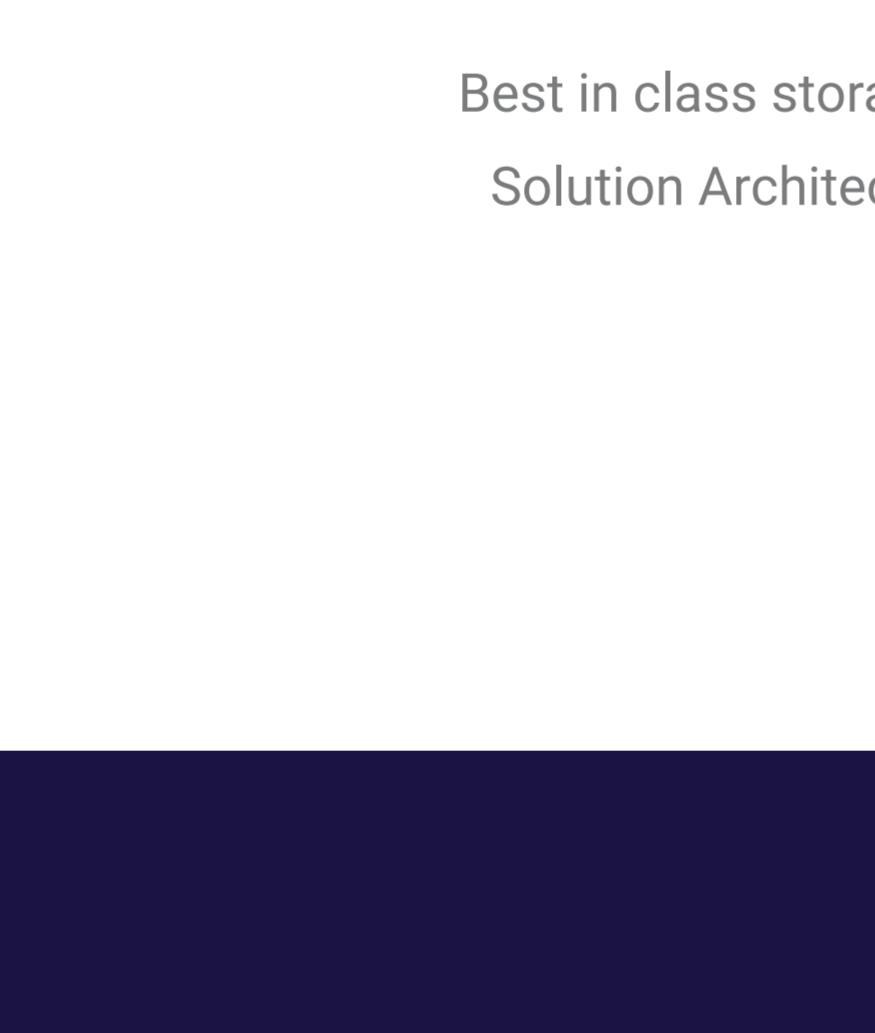
### No bottlenecks

State of the art environment producing massive amounts of IOPs over Infiniband.



### Completely self-managed

Automatic, self-served provisioning, web based system.



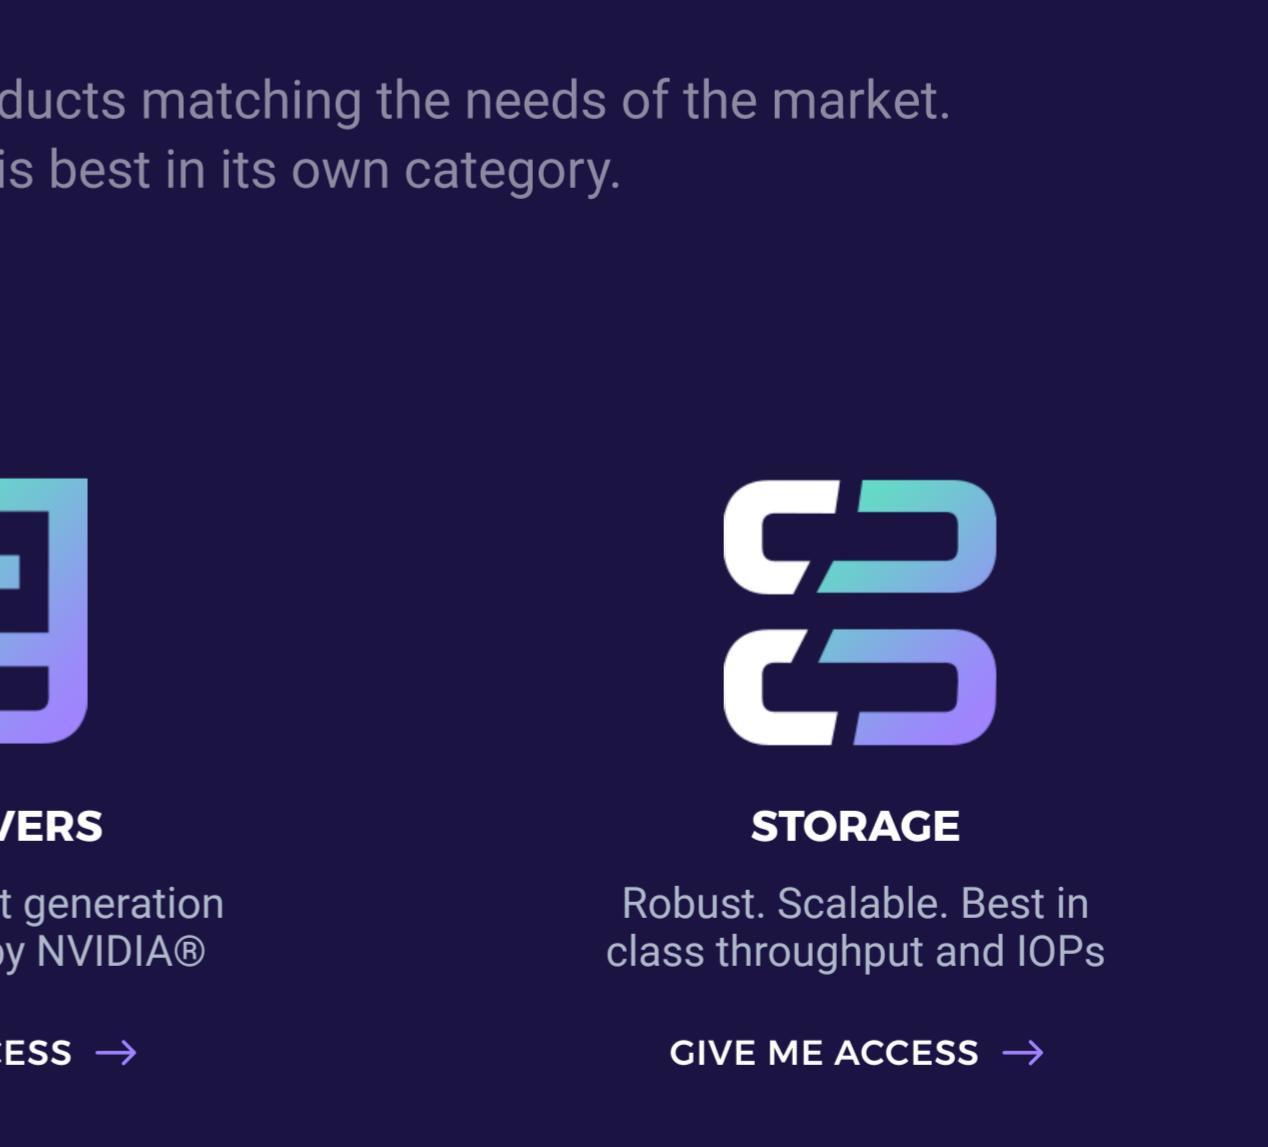
### Flexible location

Place your services in the location that will fit your needs: On-premise, your approved datacentre or our Tier-3 datacentre.



### Tailored pricing with no surprises

We customize our pricing based on your needs. You pay a flat rate for unlimited use. Predictable billing - full freedom for less money.



## Access our environment for free

Best in class storage, network and computing along with first class Solution Architects that will help you optimize your development strategy. **All for a lower cost.**



### NETWORKING

200Gb/s Infiniband.  
Lightspeed fast

GIVE ME ACCESS →



### GPU SERVERS

Multi-tiered, latest generation GPUs powered by NVIDIA®

GIVE ME ACCESS →



### STORAGE

Robust. Scalable. Best in class throughput and IOPs

GIVE ME ACCESS →

Leanfra Inc. © 2019. All rights reserved.

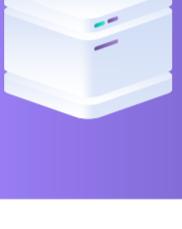
GET STARTED →

CONTACT US

**Multi-tiered, latest generation GPUs. Powered by latest Tesla®, Titan RTX® and 2080ti RTX®.**

## GPU Servers

Tiers to suit your every need

 <b>Prototyping tier</b> <span style="border: 1px solid black; border-radius: 10px; padding: 2px;">TIER #1</span> GPUs: 10x RTX 2080 Ti, Memory 11G GDDR6 Best for: and here examples of use	<b>Proof of concept tier</b> <span style="border: 1px solid black; border-radius: 10px; padding: 2px;">TIER #2</span> GPUs: 10x RTX 2080 Ti, Memory 24G GDDR6 Best for: and here examples of use	<b>Production tier</b> <span style="border: 1px solid black; border-radius: 10px; padding: 2px;">TIER #3</span> 10x PCIe Tesla V100 32G, Memory 32G HBM2 Best for: and here examples of use	<b>Lightning tier</b> <span style="border: 1px solid black; border-radius: 10px; padding: 2px;">TIER #4</span> 8x PCIe Tesla V100 32G, Memory 32G HBM2 For extremely fast communication between GPUs	<b>Absolute best tier</b> <span style="border: 1px solid black; border-radius: 10px; padding: 2px;">TIER #5</span> 16x PCIe Tesla V100 32G, Memory 32G HBM2 The absolute best on the planet
<b>GET STARTED</b> →				



**DDR4**



## Start now!

Send a request to check our environment or just simply talk with one of our experts.

**GET STARTED** →

**CONTACT US**



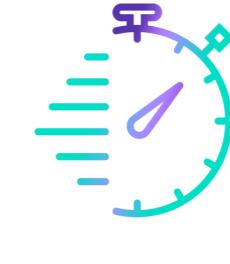
100 or 200Gb/s Infiniband. Lightspeed fast.

## Networking

Fast. Really fast.



16Tb/s of bidirectional throughput



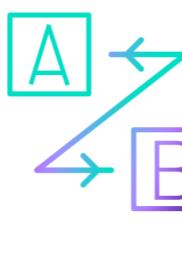
15.6 billion messages per second



90ns of port-to-port switch latency



Compliance with the Federal Information Processing Standards (FIPS)



Block-level encryption enables protection between users sharing the same resources, as various encryption keys can be used



Non-Blocking & Asynchronous ports

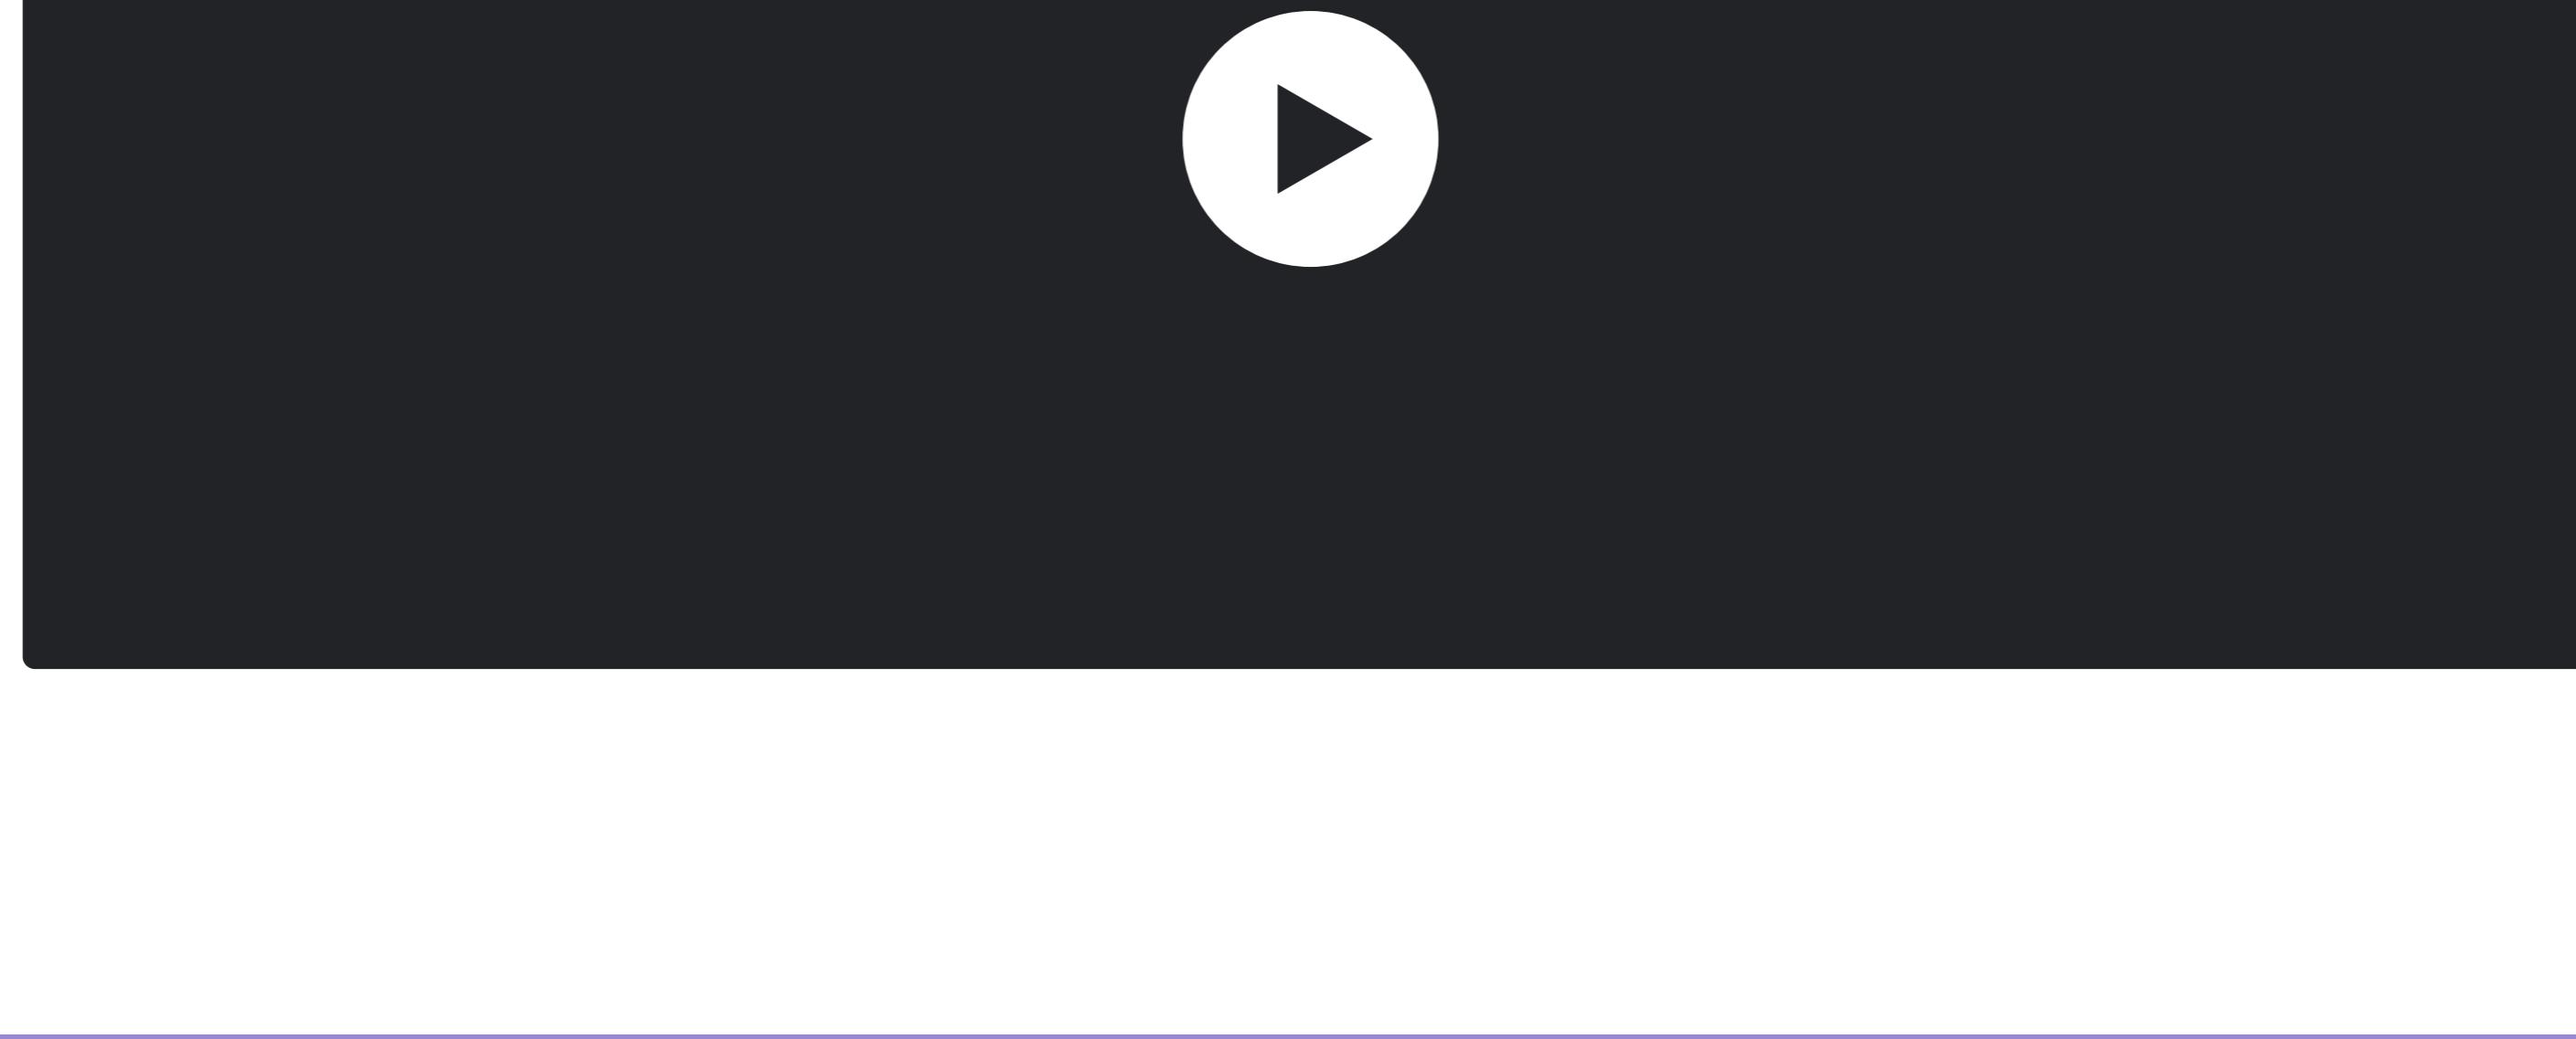
**Designed with data-centric architecture in mind: your network becomes a distributed processor.**

Enable In-Network Computing and In-Network Memory capabilities, offloading even further computation to the network, which saves CPU cycles and increases the efficiency of the network.



## Title describing this part

Subtitle, a bit longer one like 7 words



## Start now!

Send a request to check our environment or just simply talk with one of our experts.

GET STARTED →

CONTACT US

90 n/s

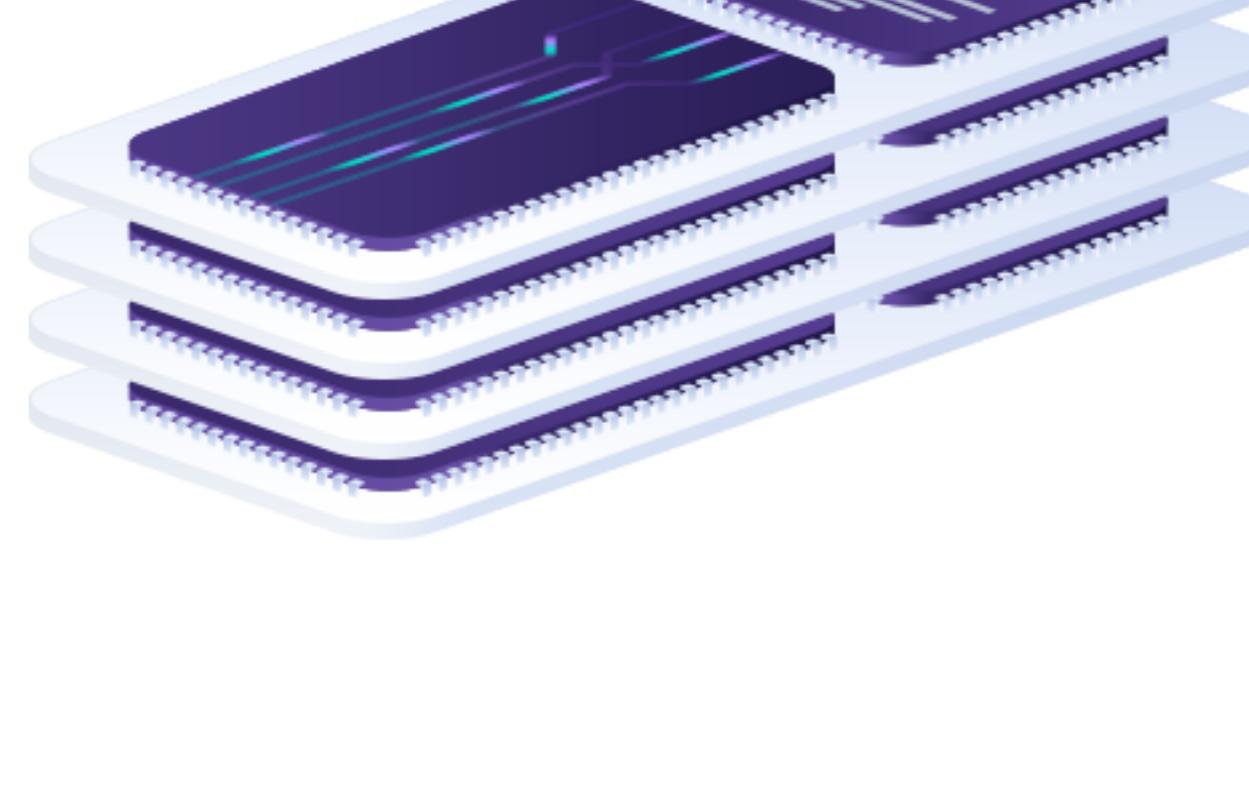
**Robust. Scalable. Infiniband based HPC architecture for extreme throughput and IOPs.**

## Storage

Cutting-edge storage.

### Full enterprise grade NVMe Storage

- Up to 732,000 IOP/s random writes
- Up to 6,708,000 IOP/s random reads
- Up to 76.8 GB/s sequential read throughput
- Up to 14.4 GB/s sequential write throughput



### 200Gb/s Infiniband network speed

- NVMe over Fabric (NVMe) Target Offloads
- Hardware offloads for VXLAN encapsulated traffic
- Sub-700 nanosecond latency



### We use ClusterFS!

- Fully POSIX compatible
- Provides replication, quotas, geo-replication, snapshots and bitrot detection.
- Scales to petabytes
- Handles thousands of clients

## Start now!

Send a request to check our environment or just simply talk with one of our experts.

[GET STARTED →](#)

[CONTACT US](#)

# The choice is yours!

We understand that servers often contain highly sensitive information and so we offer our customers as much flexibility to keep that information safe.



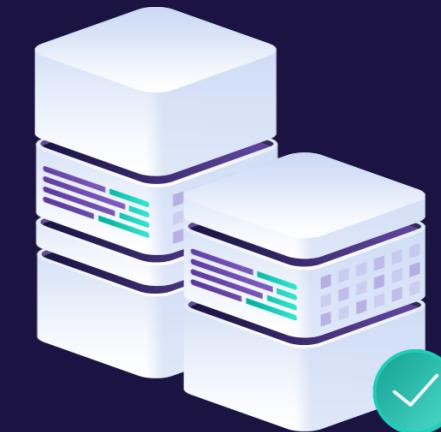
## ON-PREMISE USE OUR POD IN YOUR PREMISES

- We will ship our Pod to your premise.
- We will install and configure from A to Z.
- We will remotely maintain, monitor and support the entire Pod.

[GET STARTED →](#)

## LEANFRA'S CLOUD TIER 3 DATACENTRE

- Concurrently maintainable infrastructure
- Guaranteeing reliable service - 100% uptime
- Designed and built to provide best in class security.
- 24/7 monitoring by onsite security and Network Operations Centre personnel.
- Up to 10Gb/s WAN

[GET STARTED →](#)

## YOUR APPROVED DATACENTER USE OUR POD IN A DATACENTER

- We will ship our entire stack to your datacenter
- We will install and configure from A to Z.
- We will remotely maintain, monitor and support the entire stack.

[GET STARTED →](#)

# Tailored Pricing

Contact us and we will help to build you a quote  
that matches your requirements



## ENTERPRISE

Contact us  
for enterprise pricing

[GET STARTED →](#)



## TRIAL

Try Leanfra  
for free!

[TRY IT YOURSELF →](#)



## Try Leanfra for Free!

Try our HPC environment **for free**,  
no credit card, no commitment in 3 simple steps!

---

Step 1: Fill in your details

Step 2: Pick a date

 Full Name

 Email

 Company name

 Job title

 Phone

For approval we will need to call you for validation.

I'm not a robot



NEXT →

Step 1: Fill in your details

Step 2: Choose a date

Select a day

TODAY



TUE  
Jun 27

WED  
Jun 28

THU  
Jun 29

FRI  
Jun 30

MON  
Jul 3

TUE  
Jul 4



unavailable

Confirm?

unavailable

unavailable

THIS WEEK

NEXT WEEK



# Thank you!

for scheduling a trial in our HPC environment!

You will get an email to approve your email address and following that a phone call from us  
to confirm and apply the trial.

[GO TO BLOG →](#)

# How can we help you?

Choose a category below so we can get back to you  
as soon as possible.

**SERVICES**

Contact one of  
our Service Experts

**PRODUCT INFORMATION**

Contact Leanfra  
Tech Team Expert

**PARTNERSHIP**

I'd like to establish  
a partnership with Leanfra

**CAREER**

I'm interested in working  
with Leanfra

**MEDIA**

I have a  
media inquiry

**GENERAL INQUIRY**

I have a  
general question

# Contact our Technician

Please provide some information on how we can help you  
and we'll follow up.

 Full Name

 Email

 Company name  Job title

 Phone

 Comment

I'm not a robot 

**SEND →**

## Why Leanfra?

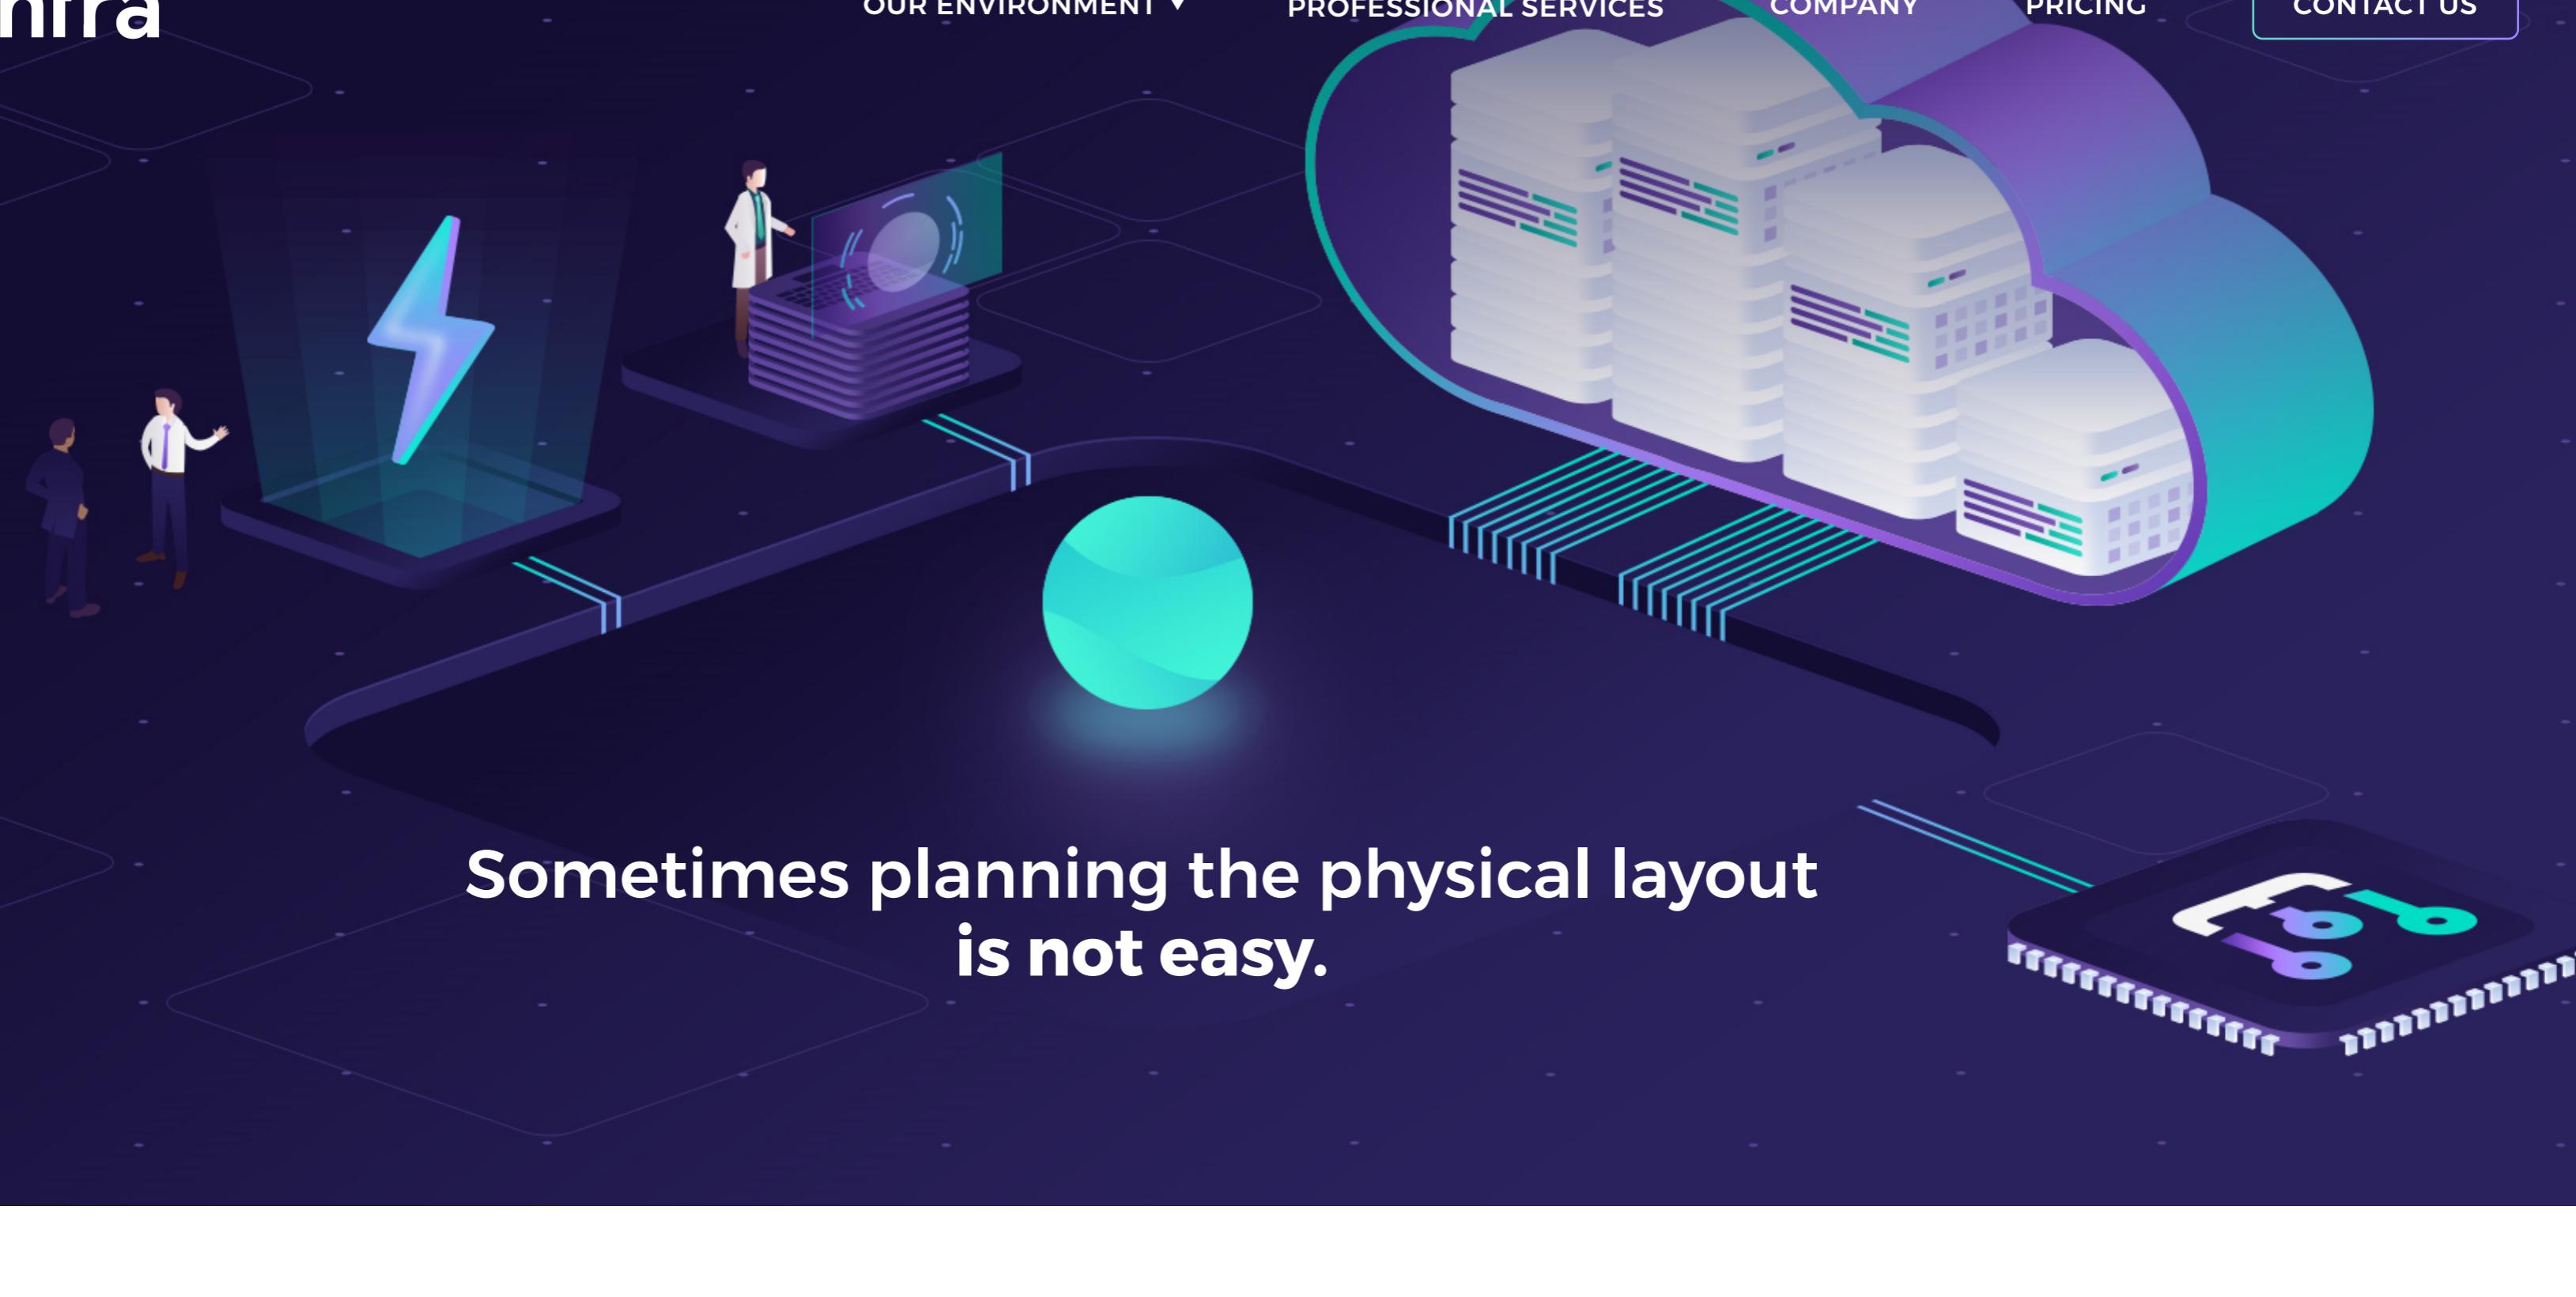
Leanfra was created to provide state-of-the-art HPC environments at an affordable price, enabling new technology to evolve at substantially lower development and research costs.

## What we do?

Leanfra exists to consistently research and benchmark for breakthrough technologies to integrate within our stack and keep you  
*[this is the text you have sent]*

## What is Our Mission?

Our goal is to keep our customers at the very cutting edge of HPC for every component in the stack for the best possible price and terms, making GPU cloud computing possible for everyone!



**Sometimes planning the physical layout  
is not easy.**

## Professional Services

Learn how our experts can help your business

**Budget vs development freedom often contradict each other.**

When developers need access to GPUs, high bandwidth storage and network resources it gets costly.

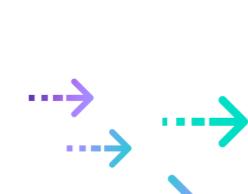


**Budget vs development freedom often contradict each other.**

When developers need access to GPUs, high bandwidth storage and network resources it gets costly.

## We can help

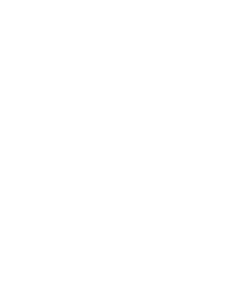
fit your product life cycle to the perfect environment  
and meet the following goals:



Implement a tiered environment



Implement more optimized and modern  
development environment.



Gives the developers 100% freedom  
to develop any time with no restrictions



Keep the budget predictable and as low as possible



Accelerate your product release in better quality



and much more...

## Start now!

Send a request to check our environment  
or just simply talk with one of our experts.

GET STARTED →

CONTACT US